

In this assignment, we will write code to compute an optimal policy for a given MDP using the algorithms that were discussed in class: Value Iteration, Howard's Policy Iteration, and Linear Programming. The first part of the assignment is to implement these algorithms. Input to these algorithms will be an MDP and the expected output is the optimal value function, along with an optimal policy. We also have to add an optional command line argument for the policy, which evaluates the value function for a given policy instead of finding the optimal policy, and returns the action and value function for each state in the same format.

MDP solvers have a variety of applications. As the second part of this assignment, we will use your solver to find an optimal policy for a batter chasing a target during the last wicket in a game of cricket.

## 1 MDP Planning

### 1.1 Problem Statement

Given an MDP, your program must compute the optimal value function  $V$  and an optimal policy  $\Pi$  by applying the algorithm that is specified through the command line. Create a python file called `planner.py` which accepts the following command-line arguments.

You are not expected to code up a solver for LP; rather, you can use available solvers as black-boxes. Your effort will be in providing the LP solver the appropriate input based on the MDP, and interpreting its output appropriately. Use the formulation presented in class. We recommend you to use the Python library PuLP

It is certain that you will face some choices while implementing your algorithms, such as in tie-breaking, handling terminal states, and so on. You are free to resolve in any reasonable way; just make sure to note your approach in your report.

### 1.2 Analysis

#### 1.2.1 Value Iteration Continuing

Value iteration is a method of computing an optimal MDP policy and its value. Value iteration starts at the end and then works backward, refining an estimate of either  $Q$  or  $V$ . There is really no end, so it uses an arbitrary end point. We define an arbitrary  $V$  and then recursively make it better. We stop the recursion when the improvement is less. Let us see the code snippet for it.

```
72 def value_iteration_cont(numStates,numActions,transition_matrix,reward_matrix,discount_factor,epsilon):
73     # Initialize Value Function arbitrarily
74     V1 = np.zeros(numStates)
75     V2 = np.zeros(numStates)
76     # Initialize optimal_policy arbitrarily
77     policy = np.zeros(numStates)
78     # Initialize differences arbitrarily
79     diff = np.zeros(numStates)
80     # Repeat until convergence
81     while True:
82         # For each state s
83         for s in range(numStates):
84             # Compute the value of each action
85             action_values = np.zeros(numActions)
86             for a in range(numActions):
87                 # Compute the value of each action
88                 action_values[a] = np.sum(transition_matrix[s,a,:] * (reward_matrix[s,a,:] + discount_factor * V1))
89             # Select the best action
90             best_action = np.argmax(action_values)
91             # Update the policy
92             policy[s] = best_action
93             # Update the value function
94             V2[s] = action_values[best_action]
95             diff[s] = np.abs(V1[s] - action_values[best_action])
96             # Check for convergence
97             if np.all(diff < epsilon):
98                 break
99             for s in range(numStates):
100                 # Update the value function
101                 V1[s] = V2[s]
102     return V1, policy
```

Value Iteration Continuing Code snippet

## 1.2.2 Policy Iteration Continuing

Policy Iteration Continuing is an algorithm to find an optimal MDP policy. Here at every step we update the existing policy using the Q values. The important part is that we never get the same policy twice. We always move in one direction. We stop when there is no scope for improvement. The policy at which we stop is the optimal policy. Let us see the code snippet for it.

```
106 #####FUNCTION TO SOLVE THE MDP USING HOWARD POLICY ITERATION FOR CONTINUOUS MDP#####
107
108 def howard_policy_iteration_cont(numStates,numActions,transition_matrix,reward_matrix,discount_factor,epsilon):
109     # Initialize Value Function arbitrarily
110     V = np.zeros(numStates)
111     # Initialize optimal policy arbitrarily
112     policy = np.zeros(numStates)
113     #Variable to store whether improvement is possible or not
114     improve=True
115     while (improve):
116         improve = False
117         # For each state s
118         #Compute the policy evaluation for current policy
119         V = policy_evaluation(policy,numStates,numActions,transition_matrix,reward_matrix,discount_factor)
120         for s in range(numStates):
121             # Compute the value of each action
122             action_values = np.zeros(numActions)
123             for a in range(numActions):
124                 action_values[a] = np.sum(transition_matrix[s,a,:] * (reward_matrix[s,a,:] + discount_factor * V))
125             # Select the best action
126             best_action = np.argmax(action_values)
127             b = best_action
128             # Check for improvement
129             if policy[s]!=best_action and action_values[b]>action_values[(int)(policy[s].item())]:
130                 improve=True
131             # Update the policy
132             policy[s] = best_action
133     return V, policy
134
135
136
```

Howard Policy Iteration Continuing Code snippet

## 1.2.3 Linear Programming Continuing

Linear Programming is a standard method to find an optimal solution to an MDP problem. Here we consider the MDP policy evaluation values as variables and write some inequalities and by using our knowledge on vector spaces we write a minimization problem on our variables. Now our Linear solver uses this constraints and provide a solution to us. If there are multiple optimal solutions, our solver gives one solution among them. Here we are using PuLP python library to do Linear Programming. Let us see the code snippet for it.

```
34 def linear_programming_cont(numStates,numActions,transition_matrix,reward_matrix,discount_factor):
35     #Initializing the LP problem
36     problem = pulp.LpProblem("MDP", pulp.LpMinimize)
37     #Initializing the variables
38     decision_variables = []
39     for s in range(numStates):
40         variable = str('v'+str(s))
41         decision_variables.append(pulp.LpVariable(variable, lowBound=0))
42     #Initializing the objective function
43     problem += (np.sum(decision_variables))
44     #Initializing the constraints
45     for s in range(numStates):
46         for a in range(numActions):
47             problem += (decision_variables[s] >= np.sum(((transition_matrix[s,a,:]*discount_factor) * ((reward_matrix[s,a,:]/
48                 discount_factor) + decision_variables))))
49     #Solving the LP problem
50     problem.solve(pulp.PULP_CBC_CMD(msg=False))
51     #Extracting the V values
52     V = np.zeros(numStates)
53     for i in range(numStates):
54         V[i]=np.float64(decision_variables[i].value())
55     #Extracting the policy
56     policy = np.zeros(numStates)
57     for s in range(numStates):
58         # Compute the value of each action
59         action_values = np.zeros(numActions)
60         for a in range(numActions):
61             action_values[a] = np.sum(transition_matrix[s,a,:] * (reward_matrix[s,a,:] + discount_factor * V))
62         best_action = np.argmax(action_values)
63         # Update the policy
64         policy[s] = best_action
65     return V, policy
```

Linear Programming Continuing Code snippet

### 1.2.4 Value Iteration Episodic

Value Iteration Episodic is method to find optimal solution of an MDP where the number of steps is finite , i.e the process stops after reaching the end states. End states are like the final destination. Here we use an algorithm similar to the one we used in Value Iteration Continuing. The policy we get finally is the optimal policy. Let us see the code snippet for it.

```
141 def value_iteration_episodic(numStates,numActions,transition_matrix,reward_matrix,discount_factor,epsilon):
142     V = np.zeros(numStates)
143     V2 = np.zeros(numStates)
144     # Initialize optimal_policy arbitrarily
145     policy = np.zeros(numStates)
146     # Initialize differences arbitrarily
147     diff = np.zeros(numStates)
148     # Repeat until convergence
149     c=0
150     while True:
151         # For each state s
152         for s in range(numStates):
153             # Compute the value of each action
154             action_values = np.zeros(numActions)
155             for a in range(numActions):
156                 action_values[a] = np.sum(transition_matrix[s,a,:] * (reward_matrix[s,a,:] + discount_factor*V))
157             # Select the best action
158             best_action = np.argmax(action_values)
159             # Update the policy
160             policy[s] = best_action
161             # Update the value function
162             V2[s] = action_values[best_action]
163             diff[s] = np.abs(V[s] - action_values[best_action])
164         # Check for convergence
165         if np.all(diff < epsilon):
166             break
167         #print(diff)
168         for s in range(numStates):
169             V[s] = V2[s]
170     return V, policy
```

Value Iteration Episodic Code snippet

### 1.2.5 Policy Iteration Episodic

Policy Iteration Episodic is an algorithm to find an optimal MDP policy. Here at every step we update the existing policy using the Q values. The important part is that we never get the same policy twice. We always move in one direction. We stop when there is no scope for improvement. The difference between Episodic and Continuing for this method is that, in Episodic we reach end states and stop there whereas in Continuing we go on forever. The policy at which we stop is the optimal policy. Let us see the code snippet for it.

```
214 #####FUNCTION TO SOLVE THE MDP USING HOWARD POLICY ITERATION FOR EPISODIC MDP#####
215 def howard_policy_iteration_episodic(numStates,numActions,transition_matrix,reward_matrix,discount_factor,epsilon):
216     # Initialize Value Function arbitrarily
217     V = np.zeros(numStates)
218     # Initialize optimal_policy arbitrarily
219     policy = np.zeros(numStates)
220     #Variable to store whether improvement is possible or not
221     improve=True
222     while (improve):
223         improve = False
224         # For each state s
225         #Compute the policy evaluation for current policy
226         V = policy_evaluation(policy,numStates,numActions,transition_matrix,reward_matrix,discount_factor)
227         for s in range(numStates):
228             # Compute the value of each action
229             action_values = np.zeros(numActions)
230             for a in range(numActions):
231                 action_values[a] = np.sum(transition_matrix[s,a,:] * (reward_matrix[s,a,:] + discount_factor * V))
232             # Select the best action
233             best_action = np.argmax(action_values)
234             b = best_action
235             # Check for improvement
236             if policy[s]!=best_action and action_values[b]>action_values[(int)(policy[s].item())]:
237                 improve=True
238             # Update the policy
239             policy[s] = best_action
240     return V, policy
```

Policy Iteration Episodic Code snippet

### 1.2.6 Linear Programming Episodic

Linear Programming is a standard method to find an optimal solution to an MPD problem. Here we consider the MDP policy evaluation values as variables and write some inequalities and by using our knowledge on vector spaces we write a minimization problem on our variables. Now our Linear solver uses this constraints and provide a solution to us. If there are multiple optimal solutions, our solver gives one solution among them.

Here we are using PuLP python library to do Linear Programming. Let us see the code snippet for it.

```
178 def linear_programming_episodic(numStates,numActions,transition_matrix,reward_matrix,discount_factor):
179     #Initializing the LP problem
180     problem = pulp.LpProblem("MDP", pulp.LpMinimize)
181     #Initializing the decision variables
182     decision_variables = []
183     for s in range(numStates):
184         variable = str('v'+str(s))
185         decision_variables.append(pulp.LpVariable(variable, lowBound=0))
186     #Initializing the objective function
187     problem += (np.sum(decision_variables))
188     #Initializing the constraints
189     for s in range(numStates):
190         for a in range(numActions):
191             problem += (decision_variables[s] >= np.sum(((transition_matrix[s,a,:]*discount_factor) * ((reward_matrix[s,a,:]/
192                 discount_factor) + decision_variables))))
193     #Solving the LP problem
194     problem.solve(pulp.PULP_CBC_CMD(msg=False))
195     #Extracting the V values
196     V = np.zeros(numStates)
197     for i in range(numStates):
198         V[i] = np.float64(decision_variables[i].value())
199     #Extracting the policy
200     policy = np.zeros(numStates)
201     for s in range(numStates):
202         # Compute the value of each action
203         action_values = np.zeros(numActions)
204         for a in range(numActions):
205             action_values[a] = np.sum(transition_matrix[s,a,:] * (reward_matrix[s,a,:] + discount_factor * V))
206         best_action = np.argmax(action_values)
207         # Update the policy
208         policy[s] = best_action
209     return V, policy
```

Linear Episodic Code snippet

## 2 Cricket: The last wicket

### 2.1 Problem Statement

The following problem is based on the game of Cricket. In Cricket, there are two teams, each with 11 players- the batting and the bowling team, which play on a pitch (between 2 wickets) located at the centre of the ground. The batting side scores runs by striking the ball bowled at the wicket with the bat and then running between the wickets, while the bowling and fielding side tries to prevent this (by preventing the ball from leaving the field, and getting the ball to either wicket) and dismiss each batter (so they are “out”). At any instance there are 2 players (striker and non-striker) from the batting team, and 11 players from the bowling team (1 bowler, 10 fielders) on the field. After every 6 balls (also called over) the striker and non-striker swap their positions: that is, the non-striker becomes a striker and vice-versa. A striking batter can either get dismissed (or out) or can score one of 0, 1, 2, 3, 4, 6 runs. To score 0, 1, 2, or 3 runs the batters need to swap their positions respective amount of times in the same ball. A striking batter can hit a 4 or 6 if the ball crosses the boundary (6 incase the ball crosses aerially).

Consider two batters A and B. In this task we aim to find the optimal policy for batter A, assuming we have no control over the actions of batter B: that is, batter A is the agent and batter B along with rest of game dynamics is part of the environment. A is a middle-order batter at the last wicket, along with a tail-end “B”. Additionally, consider that B can only either get out or score 0/1 runs. As part of the environment, we can fix a parameter “q” indicating the degree of weakness of B (a detailed description is given in the next section) The batting team still has to get T runs ( $T \leq 30$ ) in O balls ( $O \leq 15$ ). We can formulate this problem as an MDP. The set of states is encoded in the form bbr, where bb and rr are 2-digit numbers representing the number of balls left, and the number of runs to score to reach the target. Single digit numbers have a 0 attached to the left to reach 2 digits (eg - 0701 for 1 run to score in 7 balls).

### 2.2 Solution

Here the problem is divided into 3 parts. We need to first encode and then feed the output of this encoder to planner.py and then the output of planner.py to decoder.py

#### 2.2.1 encoder.py

Encoder.py python code takes the list of states provided to it and encodes it to MDP problem. Let me explain how I encoded it into MDP.

**NumStates:** Number of states in our MDP problem. In my MDP encoding  $\text{numStates} = 2\text{states} + 2$  where states are number of states provided to us in the input. Here the first  $\text{numStates}/2 - 1$  states correspond to the case where batter A is batting and the next  $\text{numStates}/2 - 1$  to the batter B. The last two states correspond to win and lose. Game ends when someone reach those states.

**NumActions:** There are five actions possible. Batsman can try to hit 0,1,2,4,6 and these corresponds to actions 0,1,2,3,4.

**Transitions:** Using the probabilities of score for a single ball for different actions given in the input file I calculated the probabilities of Transitions between states.

**End states:** Here the end states are the winning and losing states which are last 2 states.

**MDPtype:** As we have end states we solve this in episodic style.

Let us see the code snippets of it.

```
44
45 #--#OPENING THE STATES FILE
46 states = open(args.states,'r')
47 lines = states.readlines()
48 num_states = len(lines)
49 #--#TRANSITION PROBABILITIES MATRIX
50 transition_matrix = np.zeros((2*num_states+2,2*num_states+2,5))
51 #--#REWARD MATRIX
52 reward_matrix = np.zeros((2*num_states+2,2*num_states+2))
53 reward_matrix[:,2*num_states:2*num_states+1]=1
54 #--#DICTIONARY TO STORE STATE NUMBERS TO ZERO INDEXING-MAPPING
55 dict = {}
56 cnt=0
57 for line in lines:
58     arr = line.split()
59     dict[int(arr[0])] = cnt
60     cnt+=1
61 #--#WIN KEY IS SET TO 11111
62 dict[11111] = 2*num_states
63 #--#LOSS KEY IS SET TO 99999
64 dict[99999] = 2*num_states+1
```

Code Snippet To Open File And Read States From It And Initialise Matrices.

```
65 for line in lines:
66     arr = line.split()
67     num = int(arr[0])
68     rr = int(num%100)
69     bb = int((num//100)%100)
70     init_state = int(dict[num])
71     if bb>1:
72         #FOR PLAYER1 STATES
73         transition_matrix[int(init_state),int(dict[99999]),:] = p1_actions[:,0]
74         for i in range(6):
75             j = 6 if i==5 else i
76             if rr-j>0:
77                 next_state_ = dict[num-100-j] if ((j%2==0 and bb%6!=1) or (j%2==1 and bb%6==1)) else num_st
78                 transition_matrix[int(init_state),int(next_state_),:] += p1_actions[:,i+1]
79             else:
80                 next_state_ = dict[11111]
81                 transition_matrix[int(init_state),int(next_state_),:] += p1_actions[:,i+1]
82
83         #FOR PLAYER2 STATES
84         init_state2 = init_state + num_states
85         #--# OUT
86         transition_matrix[int(init_state2),dict[99999],:] += [p2,p2,p2,p2,p2]
87         if bb%6!=1:
88             #--# ZERO RUNS
89             nextstat = dict[num-100]+num_states
90             zz=(1-p2)/2
91             transition_matrix[int(init_state2),int(nextstat),:] += [zz,zz,zz,zz,zz]
92             nextstat = dict[num-101] if rr>1 else dict[11111]
93             transition_matrix[int(init_state2),int(nextstat),:] += [zz,zz,zz,zz,zz]
```

The above photo is the snippet of how we are assigning the transition probabilities based on ball remaining and runs remaining.

```

87         if bb%6!=1:
88             # ZERO RUNS
89             nextstat = dict[num-100]+num_states
90             zz=(1-p2)/2
91             transition_matrix[int(init_state2),int(nextstat),:] += [zz,zz,zz,zz,zz]
92             nextstat = dict[num-101] if rr>1 else dict[11111]
93             transition_matrix[int(init_state2),int(nextstat),:] += [zz,zz,zz,zz,zz]
94         else:
95             zz=(1-p2)/2
96             #ZERO RUNS
97             nextstat = dict[num-100]
98             transition_matrix[int(init_state2),int(nextstat),:] += [zz,zz,zz,zz,zz]
99             nextstat = dict[num-101]+num_states if rr>1 else dict[11111]
100            transition_matrix[int(init_state2),int(nextstat),:] += [zz,zz,zz,zz,zz]
101        else:
102            #FOR PLAYER1 STATES
103            transition_matrix[int(init_state),int(dict[99999]),:] += p1_actions[:,0]
104            for i in range(6):
105                j = 6 if i==5 else i
106                next_state_ = dict[99999] if rr-j>0 else dict[11111]
107                transition_matrix[int(init_state),int(next_state_),:] += p1_actions[:,i+1]
108            #FOR PLAYER2 STATES
109            #ZERO RUNS
110            nextstat = dict[99999]
111            init_state2 = init_state + num_states
112            zz = (1+p2)/2
113            transition_matrix[int(init_state2),int(nextstat),:] += [zz,zz,zz,zz,zz] #zero runs+ out probability
114            zz=(1-p2)/2
115            nextstat = dict[99999] if rr>1 else dict[11111]
116            transition_matrix[int(init_state2),int(nextstat),:] += [zz,zz,zz,zz,zz]
117

```

Code Snippet To Assign Transition Values and Rewards.

## 2.2.2 Planner.py

: Planner.py is the same code we coded in Task1. We use the default solver in task 1 to solve our MDP.

## 2.2.3 Decoder.py

: Decoder.py takes the output from planner.py and then using this and states from states file it creates the optimal Policy for our problem. Let us see the code snippet for it.

```

16
17
18 #OPENING THE STATES FILE
19 stat_file = open(args.states,'r')
20 states = stat_file.readlines()
21 stat_list = []
22 for state in states:
23     stat_list.append(state.split()[0])
24
25
26
27
28
29 #OPENING THE VALUE AND POLICY FILE
30 val_pol_file = open(args.value_policy,'r')
31 val_pol = val_pol_file.readlines()
32 cnt = 0
33 for line in val_pol:
34     arr = line.split()
35     a = int(float(arr[1]))
36     z = a if a<3 else int(a+1)
37     z = 6 if z==5 else z
38     print(str(stat_list[cnt])+' '+str(z)+' '+str(arr[0]))
39     cnt+=1
40 if cnt>=150:
41     break
42 stat_file.close()
43 val_pol_file.close()
44


```

Decoder code snippet to Decode output from planney.py.

## 2.3 Graphs

### 2.3.1 Win Probability vs B's strength(q value)


As you can see, the optimal policy probability is always more than the random policy as expected. And as the value of  $q$  increases win probability is decreasing, which is intuitively correct. As  $q$  increases the probability of B getting out increases so hence the probability of winning decreases.



Probability of Win vs B's Strength( $q$  value).

### 2.3.2 Win Probability vs Runs Required( in 10 balls)


As you can see, similar to above, the optimal policy is performing better than random policy as expected. Here the optimal policy probability of winning is strictly decreasing as expected but when you see the random policy the decrease is not consistent. At some places as runs increases the probability of winning is increasing, this behaviour is the result of random policy. As the runs increases new states gets introduced. In random policy if the current policy is a bad one and after adding some more states the policy values for new states may reduce the bad effect of old policy and hence the probability is increasing sometimes.



Probability of Win vs Required Runs(10 Balls).

### 2.3.3 Win Probability vs Balls(10 Runs)

Similar to above graphs the performance of the optimal policy is better than the random policy. Both of them have zero probability of winning when balls are 1. But the probability of winning is zero until four balls for random policy whereas it is non zero for optimal one. And one important observation is that the probability of winning is changing drastically for random policy at balls=6 and balls=12. This is because of the over change of the batsman. And also this over change is responsible for decrease of probability of winning near ball=6 and balls=12 for optimal policy.



Probability of Win vs No Of Balls(10 Runs).